

# PROJECT REPORT

---

Presented To  
SAAD MUNIR

Presented By  
HAZIQ IJAZ  
HASSAAN ULLAH

# Table Of Contents



**01** Dataset

**02** Pre  
Processing

**03** Code  
Explanation

**04** Models  
Results

**05** Website



# Dataset:

We have used three different datasets for this project

- S&P 500 dataset
- Hourly energy consumption dataset
- Daily atmosphere CO2 dataset

```
import pandas as pd

co2_data = pd.read_csv('Co2.csv')
energy_data = pd.read_csv('Energy.csv')
sp500_data = pd.read_csv('sp500_data.csv')
```

Python

## Pre Processing:

### 1. CO2 Data Preprocessing:

- Converted date columns to datetime format and set the index to 'Date'.
- Scaled CO2 concentration data using Min-Max scaling.
- Applied log transformation and differencing for stationarity analysis.

### 2. Energy Data Preprocessing

- Formatted datetime column and set it as the index.
- Interpolated missing values using time-series method.
- Filled remaining missing values with forward-fill and backward-fill.
- Scaled energy consumption metrics using Standard Scale.

### 3. S&P 500 Data Preprocessing

- Converted date column to datetime format and set it as the index.
  - Scaled stock prices and volume using Min-Max scaling.
- Applied log transformation and differencing for stationarity analysis.

## 4. Missing Values Handling

- Checked and displayed missing values for CO2, Energy, and S&P 500 datasets.
- Imputed missing values in the energy dataset through interpolation, forward-fill, and backward-fill methods.

## 5. Stationarity Analysis

- Utilized Augmented Dickey-Fuller (ADF) test to assess stationarity.
- Applied log transformation and differencing to achieve stationarity in CO2, Energy, and S&P 500 datasets.
- Displayed ADF statistic, p-value, and stationarity result for each dataset.

```
#CO2 Data
co2_data['Date'] = pd.to_datetime(co2_data[['Year', 'Month', 'Day']])
co2_data.set_index('Date', inplace=True)

#Energy
energy_data['Datetime'] = pd.to_datetime(energy_data['Date'] + ' ' + energy_data['Time'], format='%d/%m/%Y %H:%M:%S')
energy_data.drop(columns=['Date', 'Time'], inplace=True)
energy_data.set_index("Datetime", inplace=True)
for col in energy_data.columns:
    energy_data[col] = pd.to_numeric(energy_data[col], errors='coerce')

#S&P 500 Data
sp500_data['Date'] = pd.to_datetime(sp500_data['Date'], format='%d/%m/%Y')
sp500_data.set_index('Date', inplace=True)

#Missing Values
missing_values_co2 = co2_data.isnull().sum()
missing_values_energy = energy_data.isnull().sum()
missing_values_sp500 = sp500_data.isnull().sum()

# Displaying missing values
(missing_values_co2, missing_values_energy, missing_values_sp500)
```

Python

```
energy_data.interpolate(method='time', inplace=True)

# Forward fill
energy_data.fillna(method='ffill', inplace=True)

# If still missing values after forward fill, use backward fill
energy_data.fillna(method='bfill', inplace=True)

missing_values_energy = energy_data.isnull().sum()
missing_values_energy
```

Python

```

from sklearn.preprocessing import MinMaxScaler, StandardScaler

minmax_scaler = MinMaxScaler()
std_scaler = StandardScaler()

# CO2 Data
scaler = MinMaxScaler()
co2_data['CO2_ppm_scaled'] = scaler.fit_transform(co2_data[['CO2_ppm']])

# S&P 500 Data
sp500_data[['Open_scaled', 'High_scaled', 'Low_scaled', 'Close_scaled', 'Adj Close_scaled', 'Volume_scaled']] = \
    minmax_scaler.fit_transform(sp500_data[['Open', 'High', 'Low', 'Close', 'Adj Close', 'Volume']])

# Energy Data
energy_data[['Global_active_power', 'Global_reactive_power', 'Voltage', 'Global_intensity',
             'Sub_metering_1', 'Sub_metering_2', 'Sub_metering_3']] = \
    std_scaler.fit_transform(energy_data[['Global_active_power', 'Global_reactive_power', 'Voltage', 'Global_intensity',
             'Sub_metering_1', 'Sub_metering_2', 'Sub_metering_3']])

energy_data_daily = energy_data.resample('D').mean()

```

Python

Python

```

# Function to check stationarity
def test_stationarity(ts):
    adf_result = adfuller(ts)
    return {
        'ADF Statistic': adf_result[0],
        'p-value': adf_result[1],
        'Stationary': adf_result[1] < 0.05
    }

# CO2 Data - Log Transformation and Differencing
# Log Transformation and Differencing
co2_data['CO2_ppm_log'] = np.log(co2_data['CO2_ppm_scaled'] + 1)
co2_data['CO2_ppm_diff'] = co2_data['CO2_ppm_log'].diff().dropna()
co2_stationarity = test_stationarity(co2_data['CO2_ppm_diff'].dropna())

energy_data_daily['Global_active_power_log'] = np.log(energy_data_daily['Global_active_power'] + 1)
energy_data_daily['Global_active_power_diff'] = energy_data_daily['Global_active_power_log'].diff().dropna()
energy_stationarity = test_stationarity(energy_data_daily['Global_active_power_diff'].dropna())

# S&P 500 Data - Log Transformation and Differencing
sp500_data['Adj Close_log'] = np.log(sp500_data['Adj Close'] + 1)
sp500_data['Adj Close_diff'] = sp500_data['Adj Close_log'].diff().dropna()
sp500_stationarity = test_stationarity(sp500_data['Adj Close_diff'].dropna())

# Display stationarity results
(co2_stationarity, energy_stationarity, sp500_stationarity)

```

# Arima Configuration and tuning

## 1.ACF and PACF Analysis

- ACF and PACF plots were generated to identify the autocorrelation and partial autocorrelation in the differenced energy consumption data.

## 2. ARIMA Model Fitting

- ARIMA model with parameters (p=1, d=1, q=1) was fitted to the log-transformed and differenced energy consumption data.
- The model was trained on the available historical data.

## 3.Forecasting

- Utilizing the fitted ARIMA model, forecasting was performed for the next 30 days.

- Forecasted values were generated for the log-transformed energy consumption data.

#### 4. Model Evaluation


- Mean Absolute Error (MAE), Mean Squared Error (MSE), and Root Mean Squared Error (RMSE) were calculated to evaluate the forecast accuracy against the true values.
- MAE:  $\{mae\}$ , MSE:  $\{mse\}$ , RMSE:  $\{rmse\}$ .

#### 5. Forecast Visualization


- The plot illustrates the forecasted values alongside the actual values for the next 30 days.
- This visualization provides an insight into the performance of the ARIMA model in predicting energy consumption.

## Results of Arima


### CO2 dataset:



### Energy Dataset:



# S&P Dataset:



## ANN Design and Training

### 1. Data Preprocessing

- Loaded preprocessed CO2 concentration data stored as sequences.
- Created sequences of length 5 for input data and corresponding target values.

### 2. Data Splitting

- Split the dataset into training and testing sets with a 80-20 ratio for model evaluation.

### 3. Artificial Neural Network (ANN) Architecture

- Designed a feedforward ANN model with two hidden layers and dropout regularization.
- Architecture:
  - Input Layer: 64 neurons, ReLU activation function
  - Dropout Layer: 20% dropout rate
  - Hidden Layer: 32 neurons, ReLU activation function
  - Dropout Layer: 20% dropout rate
  - Output Layer: 1 neuron

### 4. Model Training

- Trained the ANN model using Adam optimizer and mean squared error loss function.
- Epochs: 50, Batch Size: 8
- Monitored validation loss to prevent overfitting.

### 5. Model Evaluation

- Evaluated the model performance on the test set.

Calculated Mean Squared Error (MSE) and Mean Absolute Error (MAE) for model assessment.



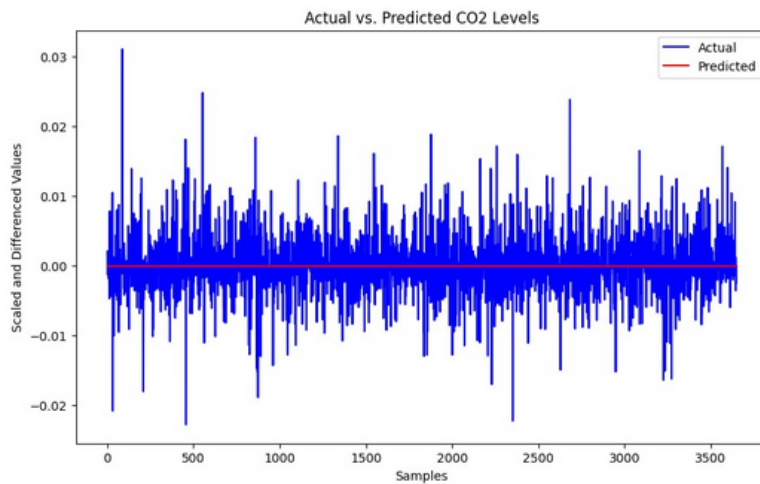
- Test Loss:  $\text{\$}\{\text{test\_loss}\}\text{\$}$ .
- MSE:  $\text{\$}\{\text{mse}\}\text{\$}$ , MAE:  $\text{\$}\{\text{mae}\}\text{\$}$ .

## 6. Results Visualization


- Visualized the actual vs. predicted CO2 concentration values.
- The plot illustrates the model's performance in capturing the temporal patterns and forecasting CO2 levels.

# Results of ANN model

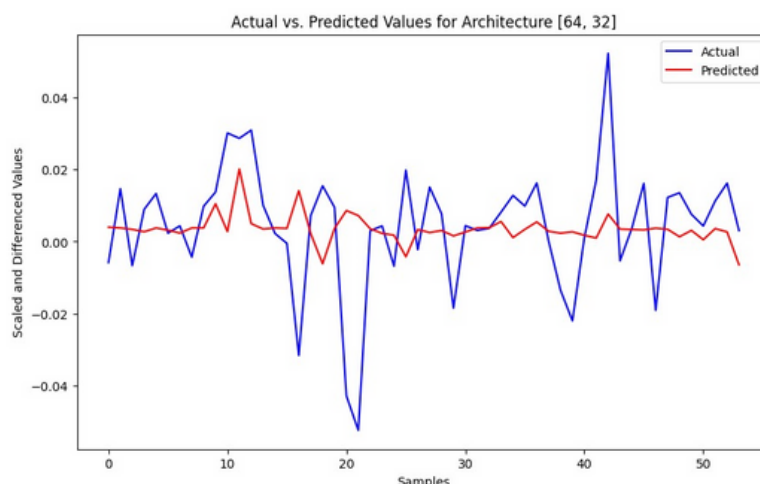
## CO2 Dataset:



## Energy Dataset:



## S&P Dataset:





# Prophet Model

## 1. Prophet Modeling

- Utilized the Prophet library for time series forecasting.
- Three datasets were prepared for forecasting: CO2 concentration, Energy consumption, and S&P 500 stock prices.

## 2. Model Preparation and Evaluation

- Each dataset was fitted to a Prophet model with appropriate settings.
- The model was trained to capture yearly and weekly seasonality patterns.
- Forecasting was performed for different time horizons: 12 months for CO2 data, 30 days for Energy data, and 12 months for S&P 500 data.

## 3. Model Evaluation Metrics


- Mean Squared Error (MSE) and Mean Absolute Error (MAE) were calculated to evaluate the forecast accuracy.
- The evaluation results provide insights into the performance of the Prophet models on each dataset.

## 4. Forecast Visualization


- Visualized the forecasted values alongside historical data using Prophet's built-in plotting functionalities.
- The plots illustrate the forecasted trends and provide a visual comparison with actual data.

# Results of Prophet model


## CO2 Dataset:



# Energy Dataset:



# S&P Dataset:



# SVR MODEL

## 1. Data Cleaning and Preparation

- Loaded CO2 concentration, Energy consumption, and S&P 500 stock price datasets.
- Cleaned data by converting date columns to datetime and handling missing values.

## 2. Support Vector Regression (SVR) Modeling

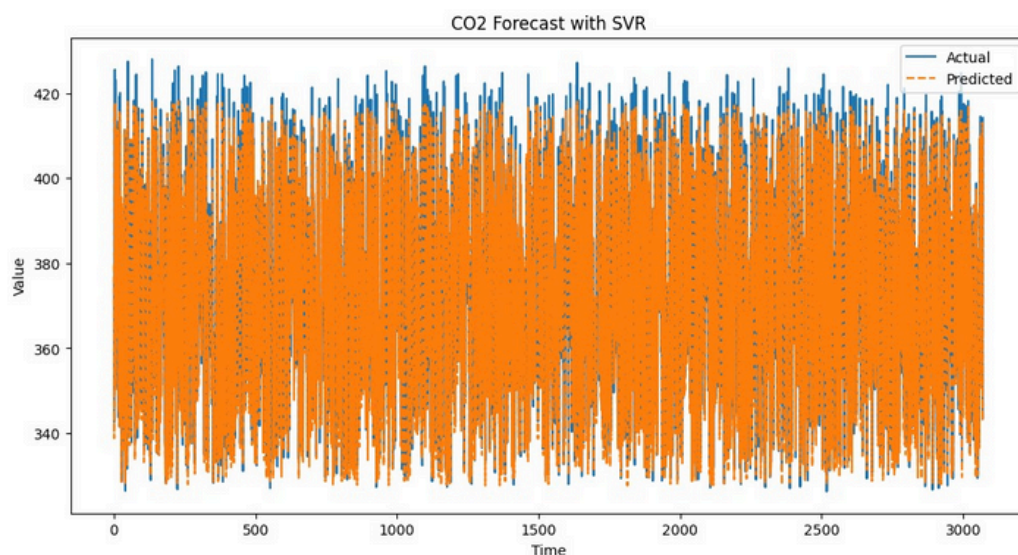
- Utilized SVR for time series forecasting.
- Implemented SVR with a linear kernel and grid search for hyperparameter tuning.
- Data were split into training and testing sets, with optional data sampling for efficiency.

### 3. Model Evaluation


- Evaluated SVR models using Mean Squared Error (MSE) and Mean Absolute Error (MAE).
- Generated plots depicting actual vs. predicted values for visual assessment.

## Results of SVR model

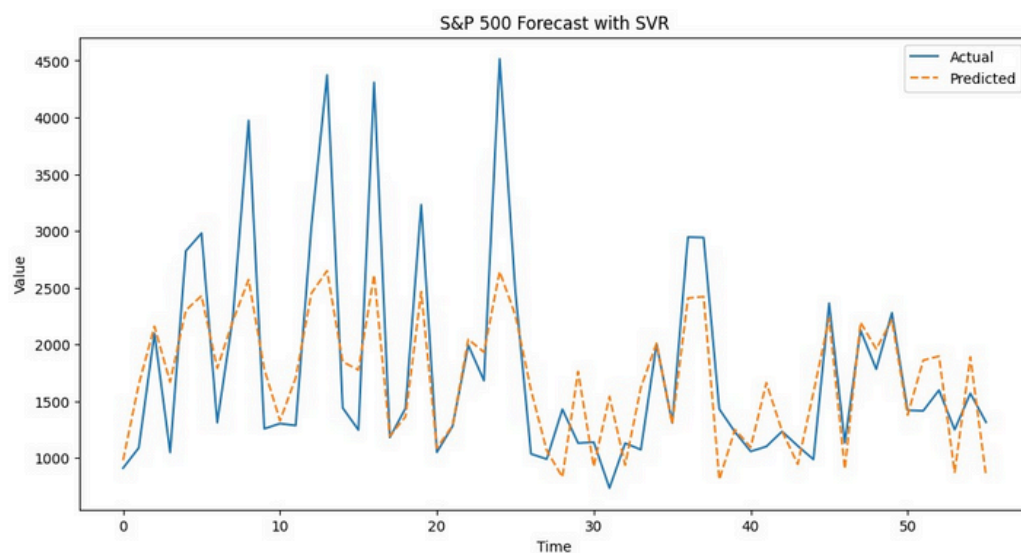
### CO2 Dataset:



### Energy Dataset:



### S&P Dataset:



# LSTM Model

## 1. Data Cleaning and Preparation

- Loaded CO2 concentration, Energy consumption, and S&P 500 stock price datasets.
- Cleaned data by converting date columns to datetime and handling missing values.

## 2. LSTM Modeling


- Utilized LSTM neural networks for time series forecasting.
- Data were prepared in sequences suitable for LSTM input.
- Implemented LSTM models with dropout regularization for better generalization.

## 3. Model Training and Evaluation

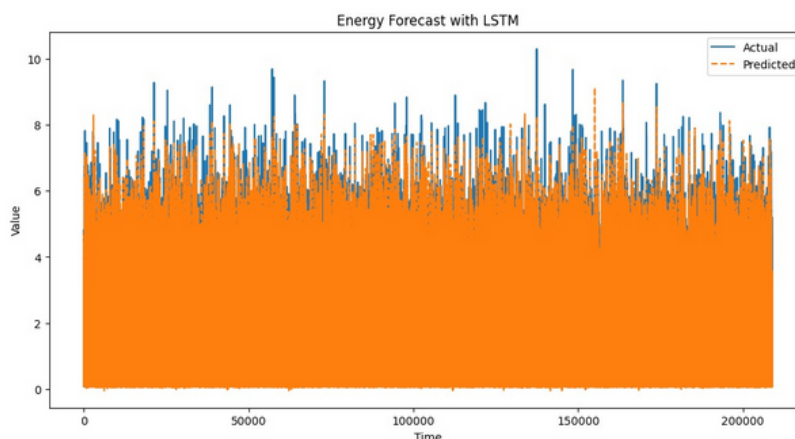
- Trained LSTM models on training data.
- Evaluated models using Mean Squared Error (MSE) and Mean Absolute Error (MAE).
- Plotted actual vs. predicted values for visual assessment.

# Results with LSTM Model

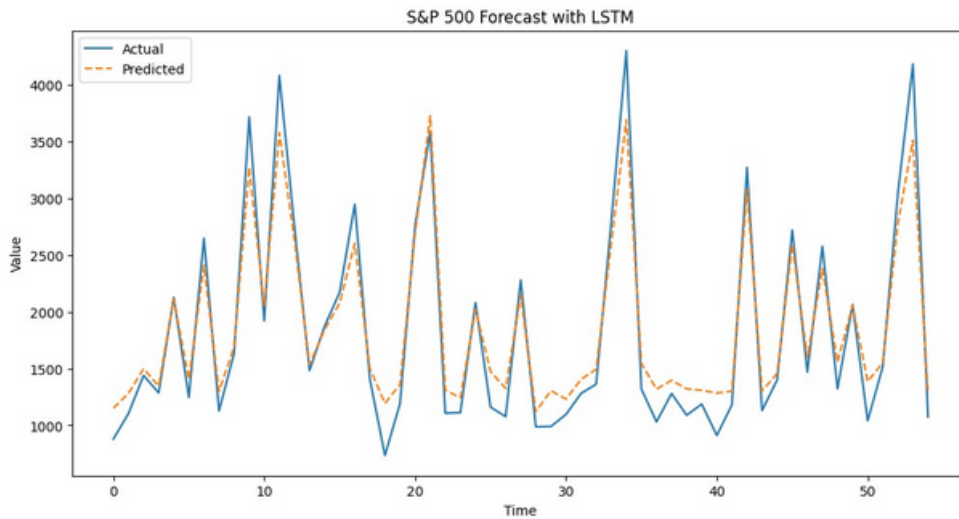
## CO2 Dataset:



## Energy Dataset:



# S&P Dataset:



## Hybrid Model

### 1. Data Preparation and Preprocessing

- Loaded CO2 concentration dataset and performed Min-Max scaling.
- Conducted log transformation and differencing to achieve stationarity.
- Prepared differenced data for ARIMA modeling.

### 2. ARIMA Modeling

- Utilized ARIMA model with order (5, 1, 0) for forecasting.
- Trained ARIMA model on differenced CO2 data.

### 3. ANN Modeling for Residuals

- Developed an Artificial Neural Network (ANN) to predict residuals from the ARIMA model.
- Scaled ARIMA predictions and residuals for training.

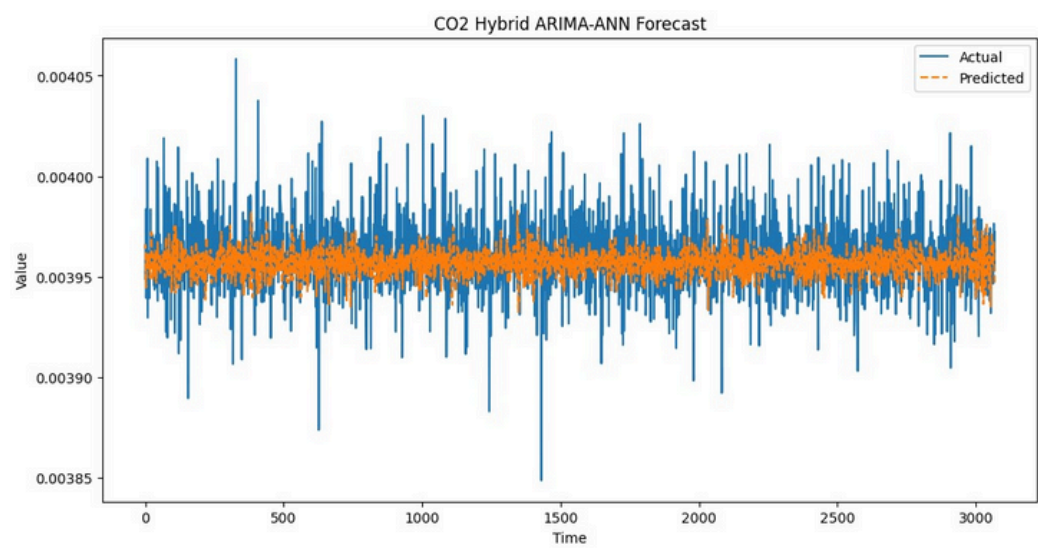
### 4. Hybrid Model Integration

- Combined ARIMA predictions with ANN predictions to obtain the final forecast.
- Evaluated the hybrid model using Mean Squared Error (MSE) and Mean Absolute Error (MAE).
- Plotted actual vs. predicted values for visual assessment.
- Similarly Hybrid model is adopted for all datasets with different models

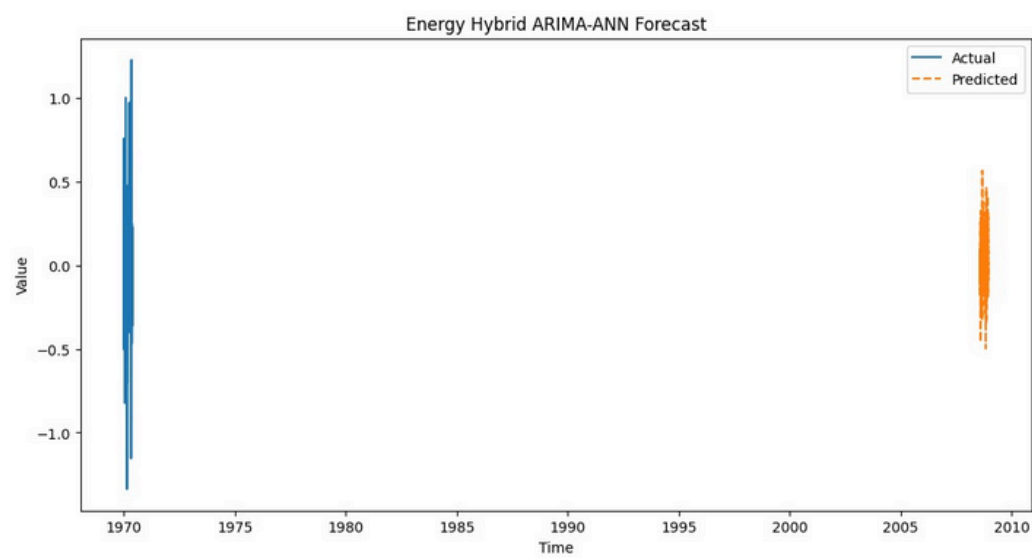
## Results with Hybrid Model



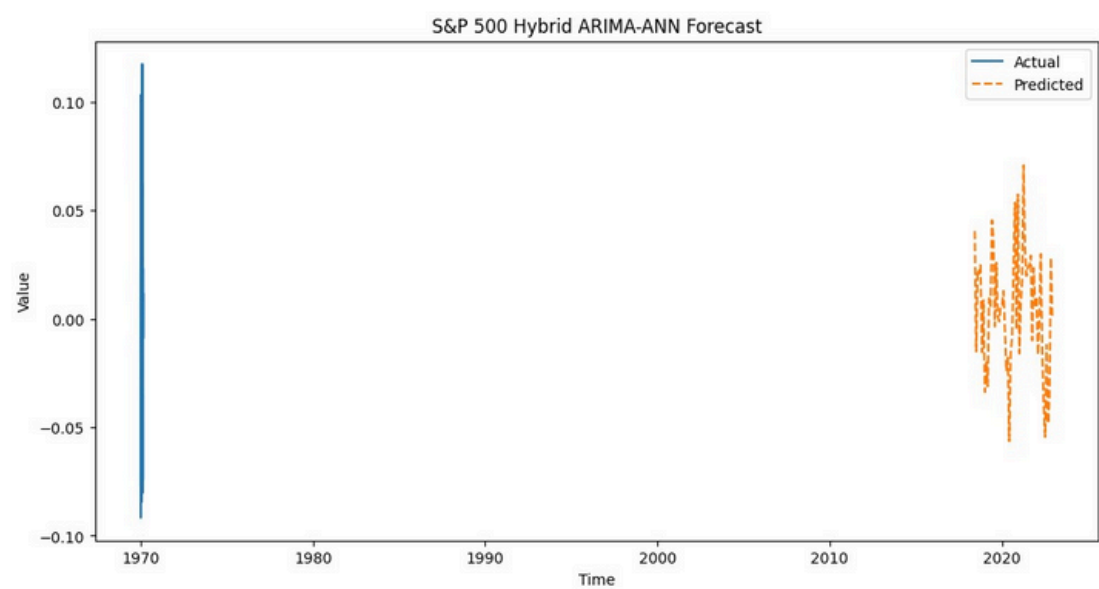
# CO2 Dataset:



# Energy Dataset:



# S&P Dataset:



# Website

